
Diverse and Faithful Knowledge-Grounded Dialogue Generation via Sequential Posterior Inference

Yan Xu^{*1} Deqian Kong^{*2} Dehong Xu² Ziwei Ji² Bo Pang³ Pascale Fung¹ Ying Nian Wu²

Abstract

The capability to generate responses with diversity and faithfulness using external knowledge is of vital importance in building a human-like trustworthy dialogue system. One line of the popular approaches is to adopt a two-step paradigm by optimizing knowledge selection and response generation separately, which may neglect the inherent correlation between these two tasks. Another line of research leverages a conditional variational model to optimize knowledge selection and response generation jointly by recruiting an inference network. In this paper, we propose an end-to-end learning framework, *Sequential Posterior Inference* (SPI), that can select knowledge and generate dialogues by approximately sampling from the posterior distribution. It does not require the inference network or assume a simple geometry of the posterior distribution. This simple and natural inference procedure of SPI can make accurate knowledge selection and generate faithful responses by directly querying the response generator. We further equip SPI with an initializer and short-run inference dynamics to explore the discrete and continuous search space efficiently and effectively. Besides modeling contributions, our experimental results on two common dialogue datasets show that SPI is superior to previous strong methods on automatic and human evaluation metrics.

1. Introduction

Open-domain dialogue systems aim at fulfilling human-machine conversations by producing human-like responses to utterances from humans (Serban et al., 2016). The emergence of large-scale pre-trained language models (PLMs) has turbocharged the development of open-domain dialogue systems (Zhang et al., 2020; Roller et al., 2021). By optimizing PLMs to maximize the token-level likelihood of gold responses given dialogue history, dialogue systems can generate fluent and natural responses. However, challenges remain to ensure that responses are diverse and informative (Ghazvininejad et al., 2018), yet remain factual and accurate (Shuster et al., 2021). Prior approaches for improving the diversity of dialogue responses focus on preventing them from being dull and repetitive (Zhao et al., 2019; Xu et al., 2022), while optimizing for diversity alone tends to encourage the dialogue system to hallucinate non-factual and inaccurate responses (Ji et al., 2022a). ChatGPT (OpenAI, 2023) tries to address this issue using a reward model trained with human preference. However, it is very resource-consuming. In order to overcome this limitation of generative dialogue systems, we need to ground system responses on external knowledge effectively. Enabling the model to access external knowledge helps prevent hallucination of non-factual content as well.

Knowledge-grounded dialogue (KGD) has been investigated in recent years (Dinan et al., 2018; Li et al., 2020; Xu et al., 2021; Yang et al., 2022). The goal is to enrich dialogue response generation to enable deep conversations without non-factual information. The task can be achieved following a two-step paradigm: (1) knowledge selection; (2) response generation. Some previous work (Lian et al.; Kim et al., 2019; Chen et al., 2020) optimizes these two steps individually. They first utilize amortized variational inference (AVI) (Kingma & Welling, 2014) for knowledge selection, where the prior distribution is conditioned on dialogue history, and the posterior distribution depends on both current response and dialogue history. Then they optimize the response generation task based on selected knowledge. Since selecting relevant knowledge based on given dialogue history is a complex one-to-many problem, it's not trivial to generate a factual response with dialogue history and se-

^{*}Equal contribution ¹Center for Artificial Intelligence Research (CAiRE), The Hong Kong University of Science and Technology, Hong Kong ²Department of Statistics, UCLA, USA ³Salesforce Research. Correspondence to: Yan Xu <yxucb@connect.ust.hk>, Deqian Kong <deqiankong@ucla.edu>.

lected knowledge solely, not to mention that inaccurate knowledge may be chosen even with a complicated knowledge selection module. Others (Liu et al., 2021) propose to bypass the knowledge selection step by providing all the knowledge candidates to the response generator, which is computationally inefficient. Therefore, it is natural to choose a probabilistic model with two latent variables to select knowledge and generate responses so that both procedures can be optimized simultaneously. CoLV (Zhan et al., 2021) follows this scheme and chooses to optimize these latent variables by recruiting an additional inference network (recognition network) to infer the posterior distribution. However, such work using AVI trained with evidence lower bound (ELBO) may ignore the fact that knowledge selection is inherently correlated to response generation and hence there might be a large amortization gap between log-likelihood and ELBO (Cremer et al., 2018). An alternative to variational inference is online posterior inference (OPI) such as Markov Chain Monte Carlo (MCMC) which may be in the form of Langevin dynamics (Langevin, 1908). (Pang et al., 2021a) proposes to generate text using short-run inference dynamics, such as finite step Langevin dynamics guided by the posterior distribution of the latent variable. OPI has demonstrated its simplicity and superiority in image modeling, trajectory prediction, etc. (Pang et al., 2020; 2021b; Xie et al., 2022; Li & Han, 2022). However, OPI-based methods are still under-explored in the scenarios of PLMs.

In this work, we propose a probabilistic model with dual latent variables, a discrete latent variable for knowledge selection, and a continuous latent variable for response generation. Instead of using AVI to approximate the posterior distribution, we propose a new approximate posterior sampling method, *Sequential Posterior Inference* (SPI). This model can be learned by approximate maximum likelihood estimation (MLE). Compared to AVI, SPI has its advantage of fewer model parameters since there is no need to parameterize the inference network, which eases the effort of fine-tuning process in PLMs. To make SPI even more efficient in PLMs, we propose to leverage the initializer and short-run MCMC to explore the search space of discrete and continuous latent variables respectively. Empirically, we show that the model trained with SPI is able to generate faithful and diverse responses with external knowledge. Our model outperforms previous methods on both WoW and Holl-E benchmarks. Further human evaluation has demonstrated its superiority as well.

Our contributions are three-fold:

- (1) We propose a probabilistic dialogue system for KGD that can be learned by approximate MLE with sequential posterior inference (SPI).
- (2) We propose to use initializer and short-run MCMC to

explore the discrete and continuous search space, which enables efficient approximate MLE learning in PLMs.

- (3) Our proposed model achieves state-of-the-art (SOTA) performance on two common KGD benchmarks.

2. Methods

2.1. Model

Suppose we have N observed examples $\{\mathcal{D}^n\}_{n=1}^N$ in dialogue dataset. For each example, $\mathcal{D}^n = (C^n, R^n)$, where C^n is the dialogue context, and R^n is the response based on the dialogue history and selected knowledge. In KGD tasks, each dialogue context consists of dialogue history H^n , and a set of M knowledge candidates $\mathbf{K}^n = \{K_i^n\}_{i=1}^M$, denoted as $C^n = (H^n, \mathbf{K}^n)$.

We consider the KGD task as a conditional generation process given dialogue history. Let $s \in \{1, \dots, M\}$ be a discrete variable indicating the choice of the knowledge candidate. Let $z \in \mathbb{R}^d$ be a d -dimensional continuous variable as a summary or abstraction of the future response, to account for the sentence-level semantics. Consider the following generative model for R ,

$$(s, z) \sim p_\alpha(s, z|C), \quad R \sim p_\beta(R|s, z, C), \quad (1)$$

where $p_\alpha(s, z|C)$ is the context-conditioned prior model parameterized by α and $p_\beta(R|s, z, C)$ is the response generation model parameterized by β .

To be specific, we may factorize the context conditional prior model as

$$s \sim p(s|C), \quad z \sim p_\alpha(z|s, C), \quad (2)$$

where $p(s|C)$ is defined as a multinomial distribution with $\mathbb{P}(s = i) = \frac{1}{M}$, $i \in \{1, \dots, M\}$ and $p_\alpha(z|s, C) = \mathcal{N}(f_\alpha(s, C), \mathbf{I})$ is an isotropic Gaussian. In our implementation, $f_\alpha(\cdot)$ is parameterized by a pre-trained BART (Lewis et al., 2020) encoder.

For the response generation model, $p_\beta(R|s, z, C)$ is defined in a conditional auto-regressive manner,

$$p_\beta(R|s, z, C) = \prod_{l=1}^L p_\beta(r_l|s, z, r_{<l}, C), \quad (3)$$

where L refers to the sentence length of the response R , r_l is the l -th token of the response, and β is parameterized by an pre-trained BART decoder. Note that s and z control every step of the auto-regressive model.

The context-conditioned distribution of response R is $p_\theta(R|C) = \sum_s \int p_\theta(s, z, R|C) dz$, where $\theta = \{\alpha, \beta\}$. Given C and R , the inference of (s, z) can be approximately

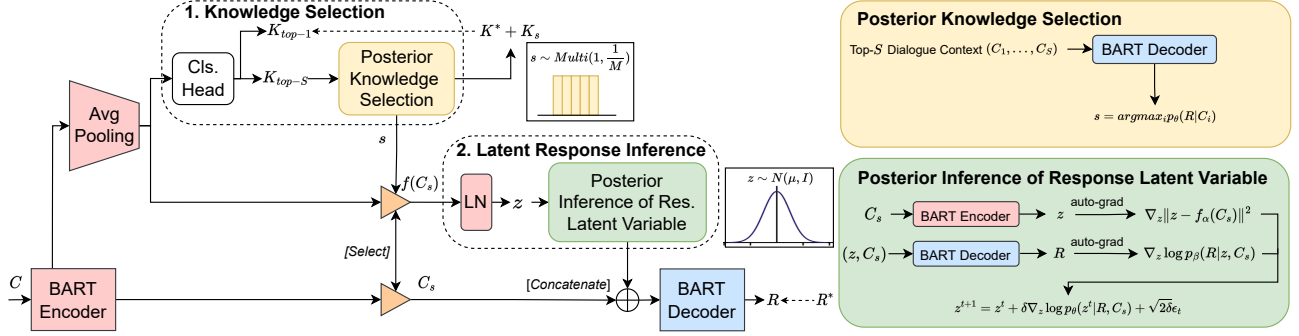


Figure 1. The overview of the learning algorithm of **SPI** (left), where modules in pink denote the context-conditioned prior model and modules in blue denote the response generation model. The prior model is instantiated with the BART encoder, while the generator is implemented with the BART decoder. The dotted lines indicate where our loss functions come from. We also demonstrate details of posterior knowledge selection and posterior inference of the response latent variable on the right.

achieved using SPI (see Section 2.3),

$$p_\theta(s, z|R, C) = \frac{p_\theta(s, z, R|C)}{p_\theta(R|C)}. \quad (4)$$

2.2. Learning

Given training examples, $\{\mathcal{D}^n = (C^n, R^n)\}_{n=1}^N$, the model can be learned using maximum likelihood estimation (MLE) where the log-likelihood is

$$L(\theta) = \frac{1}{N} \sum_{n=1}^N \log p_\theta(R^n|C^n), \quad (5)$$

where $\theta = \{\alpha, \beta\}$ are the learnable parameters of the model.

Then the gradient of the log-likelihood function can be calculated by

$$\begin{aligned} \nabla_\theta \log p_\theta(R|C) &= \frac{1}{p_\theta(R|C)} \nabla_\theta p_\theta(R|C) \\ &= \frac{1}{p_\theta(R|C)} \sum_s \int \nabla_\theta p_\theta(s, z, R|C) dz \\ &= \sum_s \int \frac{p_\theta(s, z, R|C)}{p_\theta(R|C)} \nabla_\theta \log p_\theta(s, z, R|C) dz \\ &= \mathbb{E}_{p_\theta(s, z|R, C)} [\nabla_\theta \log p_\theta(s, z, R|C)]. \end{aligned} \quad (6)$$

Although the context-conditioned distribution $p_\theta(R|C)$ is intractable due to the latent variables being integrated out, we can approximate the above expectation using Monte Carlo samples from the posterior $p_\theta(s, z|R, C)$ in Equation (6), which will be further discussed in details as sequential posterior inference in Section 2.3.

For the gradient of the log-likelihood, we have

$$\begin{aligned} \nabla_\theta \log p_\theta(s, z, R|C) &= \nabla_\theta \log p_\alpha(s, z|C) + \nabla_\theta \log p_\beta(R|s, z, C) \\ &= \nabla_\theta \log p_\alpha(z|s, C) + \nabla_\theta \log p_\beta(R|s, z, C), \end{aligned} \quad (7)$$

where the term $p(s|C)$ disappears due to its uniform probability mass.

2.3. Sequential Posterior Inference

In Equation (6), the expectation can be approximated by Monte Carlo average with samples (s, z) from $p_\theta(s, z|R, C)$. We define the pair between dialogue history H and the s -th index of knowledge in \mathbf{K} as $C_s = (H, K_s)$ with a slight abuse of notation.

We propose a sequential posterior inference method for approximate posterior inference, where we first select knowledge $p_\theta(s|R, C)$ and then infer the response latent variable $p_\theta(z|R, C_s)$.

2.3.1. POSTERIOR KNOWLEDGE SELECTION WITH TOP-S INITIALIZER

To infer the preferred knowledge index s , we shall sample from the posterior,

$$s \sim p_\theta(s|R, C) = \int p_\theta(s|z, R, C) p_\theta(z|R, C) dz. \quad (8)$$

Since the above integration is intractable, we approximate $p_\theta(z|R, C)$ in Equation (8) by a point mass at the context-conditioned prior mean. Denote $\mu = f_\alpha(C)$. Then

$$s \sim p_\theta(s|z = \mu, R, C) \propto p_\theta(s, z = \mu|R, C). \quad (9)$$

For simplicity, we will still use $p_\theta(s|R, C)$ to represent the approximate posterior distribution $p_\theta(s|z = \mu, R, C)$. For

posterior inference, we have

$$p_\theta(s|R, C) \propto p(s|C)p_\beta(R|C_s) \propto p_\beta(R|C_s). \quad (10)$$

In this case, the choice of the knowledge s is completely dependent on the response generation model. To be concrete, for each of the knowledge candidate with its history, $C_i, i \in \{1, \dots, M\}$, we first concatenate it with dialogue history H , the posterior logits are defined by

$$\mathbb{P}(s = i) = \frac{p_\beta(R|C_i)}{\sum_{i=1}^M p_\beta(R|C_i)}. \quad (11)$$

Or we can greedily choose the one that gives the best generation performance to ease the training,

$$s = \arg \max_i p_\theta(R|C_i). \quad (12)$$

However, in the case of enormous knowledge candidates (i.e. M is large), the brutal search across all M candidates might be computationally inefficient. In this case, we propose to recruit an additional linear layer, $f_\gamma(C)$, (e.g., a classification head following BART encoder) as an initializer to narrow down the search space.

Based on the output logits from $f_\gamma(C)$, we can select top- S knowledge candidates, where $S \ll M$. Then we can leverage the aforementioned process to select knowledge by sampling from the posterior,

$$\mathbb{P}(s = i) = \frac{p_\beta(R|C_i)}{\sum_{i=1}^S p_\beta(R|C_i)}. \quad (13)$$

Or greedily,

$$s = \arg \max_i p_\theta(R|C_i), i \in \{1, \dots, S\}. \quad (14)$$

This additional top- S initializer can be learned using cross-entropy loss between the predicted logits and the ground-truth label. The ground-truth label can be chosen either by gold annotations or by posterior knowledge selection, or both, which will induce the supervised, unsupervised or semi-supervised version of knowledge selection. In our experiments, we use both the gold and selected knowledge to train the initializer.

2.3.2. POSTERIOR INFERENCE OF RESPONSE LATENT VARIABLE WITH SHORT-RUN MCMC

Previous work (Rashkin et al., 2021) defines three control codes and uses them as a prefix of the inputs to indicate how the selected knowledge is presented in the gold response. It can be considered as a high-level abstraction or summary

of the future response, whereas we choose a more flexible definition of abstraction as a trainable control code or a continuous prompt that is inferred from the future response given the history and selected knowledge.

After selecting the knowledge K_s , we infer the continuous response latent variable z by sampling from $p_\theta(z|R, C_s)$

$$z \sim p_\theta(z|R, C_s) \propto p_\theta(z, s|R, C) \quad (15)$$

using Langevin dynamics

$$z^{t+1} = z^t + \delta \nabla_z \log p_\theta(z^t|R, C_s) + \sqrt{2\delta} \epsilon_t, \quad (16)$$

where $\epsilon_t \sim \mathcal{N}(0, \mathbf{I})$, t indexes the time step of the Langevin dynamics and δ is the discretization step size. The gradient term is tractable since

$$\begin{aligned} \nabla_z \log p_\theta(z|R, C_s) &= \nabla_z \log p_\theta(z, R|C_s) \\ &= \nabla_z \log p_\alpha(z|C_s) + \nabla_z \log p_\beta(R|z, C_s), \end{aligned} \quad (17)$$

where $\log p_\alpha(z|C_s) = \|z - f_\alpha(C_s)\|^2/2 + \text{constant}$ and the second term is the response generation model. Both derivatives are tractable and can be computed by back-propagation.

The Langevin dynamics in Equation. (16) involves a drift (denoted by gradient) and a diffusion term. If $z^t \sim p_\theta(z^t|R, C_s)$, the drift term $\nabla_z \log p_\theta(z^t|R, C_s)$ aims to shift the distribution of z^t towards basins of high log-posterior. $p_\theta(z|R, C_s)$ can be further recovered by smoothing with the diffusion term $\sqrt{2\delta}\epsilon_t$, which induces randomness in sampling process.

However, running sufficiently long Markov chains is computationally impractical since the back-propagation through the generator is required in each iteration according to Equation (16). Earlier works (Pang et al., 2021a) adopt short-run MCMC (Nijkamp et al., 2019) in text modeling where they propose to approximate sample from the posterior distribution with a fixed small number of steps. Here we further scale up this idea in the scenario of PLMs. That said, we propose the following sampling procedure,

$$\begin{aligned} z^0 &\sim p_\alpha(z|C_s), \\ z^{t+1} &= z^t + \delta \nabla_z \log p_\theta(z^t|R, C_s) + \sqrt{2\delta} \epsilon_t, \end{aligned} \quad (18)$$

where $t = 1, \dots, T$, and the initial state for the Markov chain is sampled from the context-conditioned prior distribution. The total length of the Markov chain is rather small (e.g. $T = 5$). Further theoretical underpinnings of this approximated sampling and learning method can be found in Appendix A.

2.4. Algorithms

Given learning iterations $\tau = 1, \dots, T_L$, the generative model with parameters $\theta = \{\alpha, \beta\}$ can be updated through

$$\begin{aligned} \theta_{\tau+1} &= \theta_{\tau} + \eta_1 \Delta\theta, \\ \Delta\theta &= \frac{1}{N} \sum_{n=1}^N \mathbb{E}_{p_{\theta_{\tau}}(s^n, z^n | R^n, C^n)} [\nabla_{\theta} \log p_{\theta}(s^n, z^n, R^n | C^n)]. \end{aligned} \quad (19)$$

The top- S initializer with parameters γ as $f_{\gamma}(C)$ can be viewed as a multi-label classifier or multiple binary classifiers and then be updated by cross-entropy loss,

$$\begin{aligned} \mathcal{L}_{\text{CE}}(y, C) &= - \sum_{i=1}^M y_i \log f_{\gamma}(C_i) + (1 - y_i) \log(1 - f_{\gamma}(C_i)), \\ \gamma_{\tau+1} &= \gamma_{\tau} - \eta_2 \frac{1}{N} \sum_{n=1}^N \nabla_{\gamma} \mathcal{L}_{\text{CE}}(y^n, C^n), \end{aligned} \quad (20)$$

where $\{y^n\}_{n=1}^N$ denotes labels which can be obtained by posterior knowledge selection and/or annotations (if we use both posterior knowledge selection and annotations, it is possible that two of y_i 's equal to 1). Therefore, by updating initializer in this way, it can output top- S candidates that likely include the gold knowledge or the one from posterior knowledge selection.

We develop a learning method that both the response generation model and the top- S initializer can be trained jointly based on SPI. The learning algorithm is summarized in Algorithm 1.

As for generation, we can get a response by greedy search as summarized in Algorithm 2. Given dialogue context $C = (H, \mathbf{K})$, we can first select the knowledge candidate with the highest logit from $f_{\gamma}(C)$,

$$s = \arg \max_i f_{\gamma}(C_i), \quad i \in \{1, \dots, M\}. \quad (21)$$

Then we use sample mean of the context-conditioned prior $p_{\alpha}(z | C_s) = \mathcal{N}(f_{\alpha}(C_s), \mathbf{I})$ to generate dialogue,

$$\hat{R} \sim p_{\theta}(R | z = f_{\alpha}(C_s), C_s). \quad (22)$$

Algorithm 1 Learning with Sequential Posterior Inference

input: Observed examples $\{C^n, R^n\}_{n=1}^{N_{\text{train}}}$, total training epochs T_L , learning rate η , number of candidates S in knowledge selection initializer, number of Langevin steps T , step size δ , initial weights θ_0, γ_0 .

output: Updated weights θ_T, γ_T .

for $\tau = 1$ **to** T **do**

1. Draw observed examples $\{C^n, R^n\}$.

2. **Sequential Posterior Inference**

2.1 *Posterior Knowledge Selection*

(a) Select knowledge candidates by top- S initializer $f_{\gamma}(C^n)$.

(b) Infer s from the top- S candidates using (14).

2.2 *Posterior Inference of Response Latent Variable*

Infer z by T -step short-run MCMC (18) with step size δ .

3. **Update Model Parameters**

Update θ and γ according to Equations (19) and (20).

end for

Algorithm 2 Knowledge-Grounded Response Generation

input: Observed examples $\{C^n\}_{n=1}^{N_{\text{test}}}$.

output: Response $\{\hat{R}^n\}_{n=1}^{N_{\text{test}}}$.

for $n = 1$ **to** N_{test} **do**

1. Draw test example C^n .

2. Select s using $f_{\gamma}(C^n)$ according to (21).

3. Set z as the sample mean of the prior $p_{\alpha}(z | C_s^n)$.

4. Generate \hat{R}^n by decoder using (22).

end for

3. Experiments

3.1. Experiment Settings

Datasets We conduct our experiments on two KGD datasets, Wizard of Wikipedia (WoW) (Dinan et al., 2019) and Holle (Moghe et al., 2018). In WoW, dialogues are directly grounded on the knowledge sentences retrieved from Wikipedia. 22.3k dialogues with 202k turns in WoW dataset are divided into training, validation, and test subsets. Both validation and test sets consist of seen and unseen sets, where the unseen set consists of the dialogues with the unseen initial topics during the training time. For a balance between task learning and generalizability, we merge two validation sets to select the best checkpoint. The Holle dataset contains 9k conversations with 90k utterances about movies. Each response is obtained based on unstructured knowledge such as plots, comments, and reviews about the movie. In both datasets, the gold label for knowledge selection is provided along with each dialogue turn.

MODEL	WoW SEEN								WoW UNSEEN							
	PPL↓	B3↑	B4↑	R1↑	R2↑	DIST-1↑	DIST-2↑	ACC↑	PPL↓	B3↑	B4↑	R1↑	R2↑	DIST-1↑	DIST-2↑	ACC↑
BART _{SKT}	20.3	7.6	4.4	19.4	5.4	6.8	30.3	26.8	22.3	—	4.6	19	4.7	5.2	24.5	18.3
BART _{FID}	9.5	7.9	5.8	20.9	7.8	10.4	39.6	—	10.5	8.1	6.1	20.9	7.9	6.7	24.2	—
ZRKG	40.4	2.8	1.8	18.6	2.4	5.4	22.5	—	41.5	18.6	1.1	18.5	2.4	3.4	15.6	—
DRD	23	7.5	5.5	18.0	—	—	—	—	25.6	16.5	4.3	16.5	—	—	—	—
PIPM	42.7	—	3.3	19.9	7.3	—	26.4	27.7	65.7	—	2.5	17.6	5.4	—	17.7	19.4
CoLV	39.6	—	2.9	20.6	7.9	—	29.7	30.1	54.3	—	2.1	19.7	6.3	—	20.1	18.9
KAT-TSLF	14.4	9.1	6.7	21.7	7.6	9.5	38.3	—	15.8	8.3	6.0	20.7	7.2	6.7	26.0	—
KNOWLEDGPT	19.2	9.5	7.2	22.0	7.9	8.9	36.2	28.0	22.3	8.3	6.0	20.5	6.7	6.0	23.8	24.0
SPI(OURS)	17.1	10.2	7.7	22.7	8.8	10.8	40.9	36.2	19.1	9.6	7.3	22.0	8.5	6.9	24.3	34.6
1/2 DATA	18.2	9.7	7.3	21.8	8.1	10.6	40.6	34.3	20.1	9.2	6.9	21.1	7.7	6.5	23.0	33.2
1/4 DATA	18.7	9.3	6.9	21.6	7.8	10.1	39.0	33.6	20.7	8.9	6.6	20.9	7.3	6.3	23.1	32.5
1/8 DATA	20.3	7.9	5.7	20.2	6.7	9.4	35.8	31.4	22.0	8.1	6.0	19.6	6.5	5.8	20.7	30.6
1/16 DATA	22.0	7.0	4.9	18.7	5.6	8.9	34.0	27.5	23.6	7.2	5.2	18.5	5.7	5.7	20.8	27.0

Table 1. Automatic evaluation results on WoW test sets. PPL is short for Perplexity; B3 and B4 represent BLEU-3 and BLEU-4; R1 and R2 denote Rouge-1 and Rouge-2; Dist-1 and Dist-2 denote uni-gram and bi-gram distinct metrics. Numbers of previous models are taken from (Zhao et al., 2019; Li et al., 2020; Chen et al., 2020; Zhan et al., 2021; Zhao et al., 2020; Liu et al., 2021). Our proposed method achieve new SOTA performance on WoW tet sets. The performance of our proposed model under the low-resource settings is shown in the last four rows.

MODEL	ORACLE PERFORMANCE					FEQA	QUESTEval	
	PPL↓	B3	B4	R1	R2		RD	RF
<i>WoW Seen</i>								
KNOWLEDGPT	9.1	19.2	15.5	34.5	17.3	48.1	42.2	43.5
SPI (OURS)	8.7	20.0	16.3	36.1	18.7	49.2	44.4	46.0
<i>WoW Unseen</i>								
KNOWLEDGPT	9.8	18.3	14.6	33.8	16.5	47.4	41.0	42.2
SPI (OURS)	9.2	20.1	16.3	36.0	18.7	49.6	44.0	45.7

Table 2. The results on automatic faithfulness metrics on WoW test sets. The proposed model, SPI, consistently outperforms KnowledGPT on all the metrics, showing its superior faithfulness.

Implementation Details An overview of the model structure of SPI is illustrated in Figure 1. Our training implementation is based on pre-trained BART-base (Lewis et al., 2020). The prior model (parameterized by α) and the generator (parameterized by β) are highlighted in pink and blue, respectively. The prior model is instantiated with the BART encoder, while the generator is implemented with the BART decoder. The inferred response latent variable z is concatenated with the representation of dialogue context C_s produced by the BART encoder on the dimension of sequence length, acting as a special token or trainable control code. BART decoder generates responses conditioned on z through the cross-attention mechanism in each Transformer layer. More details about the loss functions that we leverage to train the model are included in Appendix B.

Training Details We train our model with Adam optimizer with a learning rate of $1e-7$ and a weight decay of 0.005. A linear scheduler is utilized to adjust the learning rate for each step. The batch size is set as 32. We train our model on NVIDIA Geforce A6000 GPU with 15 epochs and select the best checkpoint with the lowest loss on the

validation set as our final model. The responses are generated using greedy search. We set S as 5 for knowledge selection initialization. For Langevin dynamics, the number of Langevin steps and step size are 5 and 0.1, respectively. We discuss the training time cost with Langevin dynamics in Section 3.4.

3.2. Evaluation

Automatic Evaluation To evaluate the knowledge selection performance of our model and baselines on WoW and Holl-E datasets, we use the accuracy (Acc) score, which is the ratio of the test samples where selected knowledge candidates are exactly the same as the gold annotations. As for estimating the quality of generated responses from different models, we utilize the classical overlap-based metrics: BLEU-3 (B3), BLEU-4 (B4) (Papineni et al., 2002), Rouge-1 (R1) and Rouge-2 (R2) (Lin, 2004) to measure the distance from the golden answers. Perplexity (PPL) is the exponential negative log-likelihood of the model generating gold responses. We use distinct scores (Dist-1 and Dist-2) (Li et al., 2016) to calculate the ratio of distinct uni-gram and bi-grams at the corpus level, which reflect the diversity of generated responses.

Moreover, we adopt automatic metrics, especially for evaluating the faithfulness of the generated responses, including FeQA (Durmus et al., 2020), QuestEval (Scialom et al., 2021), and the overlap-based performance given the oracle knowledge. FeQA and QuestEval are both question-answering-based frameworks for evaluating the faithfulness of the generations. They rely on iterations of question generation based on the generated text and question answering (QA) given the context. The QA performance’s accuracy is considered equivalent to the degree of faithfulness. QuestE-

MODEL	PPL↓	B4	R1	R2	DIST-2	ACC
SKT	48.9	-	29.8	23.1	-	29.2
DUKENET	42.7	19.2	32.6	19.6	28.5	30.4
PIPM	39.2	18.3	30.8	24.0	27.2	30.7
CoLV	34.8	20.3	32.0	25.8	29.9	32.7
SPI (OURS)	12.6	30.7	38.3	31.7	30.6	38.3

Table 3. Automatic evaluation results on Holl-E test set. Numbers of previous models are taken from (Kim et al., 2019; Meng et al., 2020; Chen et al., 2020; Zhan et al., 2021). Our model outperforms all the strong baselines and achieves new SOTA performance.

val has two modes: (1) reference-dependent (RD) mode assesses the generated text with one or multiple ground-truth references, and (2) reference-free (RF) mode conducts the assessment when no gold reference is available.

Human Evaluation For a comprehensive evaluation, we use human evaluation to compare the generated responses from our model with those from one of the previous SOTA models, KnowledGPT (Zhao et al., 2020).¹ We assess the responses quality from three aspects: *Fluency*, *Relevance*, and *Faithfulness*. *Fluency* assesses whether the response is complete, grammatically correct, and self-consistent without repetition, while *Relevance* evaluates whether the selected knowledge and the corresponding response are relevant to the dialogue history. Both fluency and relevance are assessed using A/B testing. We evaluate *Faithfulness* using a 4-point Likert scale. The annotators are required to compare the response with the selected knowledge, and dialogue history to decide whether the response is faithful or not. A faithful response should be fully supported by the dialogue context consisting of external knowledge and history and correctly convey the information in external knowledge. 50 data samples are randomly selected from each test set, and we ensure that each sample is evaluated by three annotators. Further details and annotator instructions are included in Appendix E.

3.3. Results

Table 1 and Table 3 report automatic evaluation results of our proposed model on WoW and Holl-E test sets. We compare our model with a number of previous strong models on both datasets and highlight the best performance of each metric in bold. The baseline models are introduced in Appendix C. Our proposed method achieves new SOTA performance on both datasets. It outperforms all the previous strong baseline models on knowledge selection accuracy and overlap-based metrics, indicating a higher quality of knowledge selection and response generation. For WoW test sets, compared with

¹Human evaluation is conducted on Amazon Mechanical Turk (<https://www.mturk.com/>).

MODEL	FLUENCY		RELEVANCE		FAITHFULNESS	
	SEEN	UN.	SEEN	UN.	SEEN	UN.
KNOWLEDGPT	62.5%	60.3%	70.8%	62.2%	3.33	3.42
SPI (OURS)	88.7%	83.3%	79.8%	74.4%	3.66	3.65

Table 4. Human evaluation results on WoW test sets, in terms of *Fluency*, *Relevance*, and *Faithfulness*. *Un.* is short for the unseen set. A pairwise t-test is conducted to validate the significance of the improvements, and the corresponding results in bold are significantly better than those from the baseline model ($p < 0.05$).

one of the previous SOTA models, KnowledGPT, our model shows a 6.9% on BLEU-4, 11.4% on Rouge-2, and 29.3% on accuracy on the test seen set. Meanwhile, the improvements on the test unseen set of WoW are even larger (e.g., 21.7% on BLEU-4, 26.9% on Rouge-2, 44.2% on accuracy). This proves the better *generalizability* of our model. Needless to say, our model pushes the SOTA performance by a large margin. The improvements on the Holl-E dataset are at least 17% for all the metrics except Distinct-2 (2%).

Furthermore, the proposed model consistently outperforms KnowledGPT on all the automatic faithfulness metrics in Table 2, showing its superior faithfulness. Our advantage over other models in distinct scores (Table 1 and Table 3) also shows that our model tends to generate more diverse responses, especially in the seen domain. In WoW unseen set, SPI outperforms KAT-TSLF (Liu et al., 2021) on Distinct-1, whereas KAT-TSLF performs better on the Distinct-2 metric. KAT-TSLF proposes a BART-based model pre-trained on a large dialogue corpus with pseudo-knowledge pairs and then adapted to WoW dataset through fine-tuning. Regarding its performance on other metrics, we believe the pre-training process is the major contributor to diversity. Our model achieves the second-best performance on Distinct-2 with no additional pre-training step or data resource.

PPL scores of our model are less satisfying than the deterministic models, i.e. BART-based FiD (Izacard & Grave, 2021). However, it is necessary to emphasize that even though there is a correlation between PPL and human evaluation to some extent, it is not directly reflecting the quality of response generation when the PPL is low, because of the likelihood trap confirmed in (Zhang et al., 2021).²

Table 4 lists the human evaluation results on both test sets of the WoW dataset, comparing KnowledGPT and our model in terms of *Fluency*, *Relevance*, and *Faithfulness*. The details about how each score is calculated are stated in Appendix E. A pairwise individual t-test validates the significance of the advantages of our model over KnowledGPT. The results in bold denote the corresponding improvements are statistically significant ($p < 0.05$). Our model is more

²If the PPL of the model is too low, the correlation with human judgment decreases.

TOP-S	WoW SEEN						WoW UNSEEN					
	B-4	R-2	DIST-2	FEQA	Q.E.(RD/RF)	ACC	B-4	R-2	DIST-2	FEQA	Q.E.(RD/RF)	ACC
1	7.3	8.4	36.6	40.4	41.1/43.0	37.0	6.9	7.7	22.5	39.2	39.9/41.8	34.7
3	7.4	8.3	39.4	40.7	41.4/43.2	34.1	7.0	7.8	22.5	40.5	40.5/42.2	32.2
5 (OURS)	7.7	8.8	40.9	49.2	44.4/46.0	36.2	7.3	8.5	24.3	49.6	44.0/45.7	34.6
10	7.2	8.8	41.1	48.0	42.4/44.2	36.4	7.3	8.4	24.4	47.7	42.3/44.0	34.6

Table 5. Ablation study on the impact of the choice of top- S for posterior knowledge selection initialization on WoW test sets. Q.E. is short for QUESTEVAL. Our final model with top-5 knowledge candidates shows a balance between diversity and overlap-based accuracy on the quality of generated responses.

LANGEVIN STEPS	WoW SEEN					WoW UNSEEN					Tr. TIME (/EPOCH)
	B4	R2	DIST-2	FEQA	Q.E.(RD/RF)	B4	R2	DIST-2	FEQA	Q.E.(RD/RF)	
0	7.4	8.7	40.3	47.4	43.8/45.6	6.9	8.2	23.5	48.0	42.9/44.6	3.50HRS
1	7.6	8.7	40.3	47.9	44.2/45.9	7.4	8.4	23.1	47.9	43.5/45.1	3.56HRS
5 (OURS)	7.7	8.8	40.9	49.2	44.4/46.0	7.3	8.5	24.3	49.6	44.0/45.7	3.68HRS

Table 6. Ablation study on the impact of the number of Langevin steps on WoW test sets. Q.E. is short for QUESTEVAL. We also present the training time (Tr. Time) per epoch under each setting. As the number of Langevin steps increases, the performance on the test seen set consistently improves, while the training time cost also increases slightly.

likely to generate fluent responses compared with Knowl-edGPT. The advantage in *Relevance* also proves that our model is capable of selecting more relevant knowledge and ensuring coherence to the dialogue history, especially in the unseen domain. According to the criteria of Faithfulness evaluation, both Knowl-edGPT and our model generate partially faithful responses. Human-like dialogue skills encourage dialogue models to respond to the users with new information that does not appear in the dialogue history. However, a diverse response may have the probability of being non-factual. Compared with Knowl-edGPT, our model generates significantly more faithful responses, while enhancing diversity given the Distinct scores in Table 1. Moreover, a case study is also included in Appendix D.

3.4. Ablation Study

Low-resource settings Our model demonstrates high training efficiency under low-resource settings. We train our model using the same hyper-parameter settings with $1/2$, $1/4$, $1/8$, and $1/16$ number of data samples on WoW datasets. From Table 1, with the increasing number of data samples during training, the performance of all the metrics improves consistently. With only a quarter of the data samples, our model is still able to perform comparably or even better than that of other strong baseline models. We compare our performance with KAT-TSLF under low-resource settings, as shown in Table 7. Our model with SPI appears to drop less on the performance under $1/4$ and $1/8$ data settings, using much less training cost than KAT-TSLF. KAT-TSLF relies on pre-training with a large dialogue corpus to prevent the model from poor performance under the low-resource setting. Because of pre-training, KAT-TSLF shows zero-shot

KGD ability and gets better diversity in some low-resource settings. However, we barely find the difficulty of applying SPI for pre-training with a large dialogue corpus.

Impact of top- S selection To study the impact of the choice of S in posterior knowledge selection with top- S knowledge candidates, we conduct experiments when $S = 1/3/5/10$ with all the other settings kept the same. As the results listed in Table 5, when the initializer is only optimized on gold labels for knowledge selection without posterior knowledge selection ($S = 1$), the model performs the best knowledge selection accuracy. However, with more and more knowledge candidates produced by the initializer, the diversity of generated responses is on the rise, whereas the best overlap-based accuracy achieves with top-5 knowledge candidates. It shows that injecting posterior information into the initializer during training may not improve knowledge selection but improves the faithfulness (FeQA and QuestEval scores) of the generated responses. This verifies our assumption about the inherent correlation between knowledge selection and response generation. It also proves that better knowledge selection helps with better results but does not guarantee better responses because the generation can still hallucinate and deviate from the knowledge source provided. The two paradigms of KGD tasks should be optimized jointly.

Impact of the number of Langevin steps In Table 6, we further study the impact of the number of Langevin steps on response generation. When training these models, all the experimental settings except the number of Langevin steps are kept the same as those in Section 3.1. When no Langevin step is taken, the response latent variable z degenerates to

be a deterministic representation. Posterior inference of z further boosts the performance of the SPI model on overlap-based accuracies, demonstrating the effectiveness of the proposed method, especially in the unseen domain. It also improves both diversity and faithfulness by providing a high-level abstraction of the further response with the response latent variable. Posterior inference with Langevin dynamics requires the model to use MCMC, which sequentially queries the BART decoder to obtain the gradient from the generator for updating the response latent variable z . One possible concern is the increasing training cost when more Langevin steps are taken. We calculate the training time per epoch for models with different Langevin steps. Posterior inference of response latent variable z with Langevin steps as five only extends the training time per epoch by 5.1%, which does not bring much burden on the training process.

4. Related Work

Knowledge-Grounded Dialogue Generation KGD task has been investigated for many years (Dinan et al., 2019; Feng et al., 2021). Existing methods pay efforts to tackle the challenges of knowledge selection and knowledge-aware response generation. Due to one-to-many problems in knowledge selection, one line of existing work adopts AVI-based methods, which construct a latent variable for knowledge selection and optimize it with variational inference (Lian et al.; Kim et al., 2019; Li et al., 2020; Chen et al., 2020). Further explorations extend the formulation to two collaborative latent variables to further augment response generation or enhance knowledge selection. (Zhan et al., 2021) utilizes two collaborative latent variables to model the distributions of knowledge and response simultaneously, while (Fu et al., 2022) proposes a variational method and introduces two latent variables to indicate the fragment of personal memory to evoke and the knowledge candidate to select respectively. Another line of research bypasses the knowledge selection step but relies on improving knowledge usage during knowledge-aware response generation given all the knowledge sentences (Zhao et al., 2020; Liu et al., 2021).

Since PLMs hallucination problem (Ji et al., 2022a) leads to some of the challenges in faithfulness, we note that to reduce hallucination in KGD systems, existing work focuses on guiding the model on correct knowledge usage (Rashkin et al., 2021; Ji et al., 2022b) or improving knowledge selection performance, thus providing dialogue models with better knowledge augmentation (Shuster et al., 2021). In this work, SPI jointly improves both processes and shows a significant faithfulness advantage through automatic and human evaluation.

Posterior Inference (Han et al., 2017) propose to learn generative image modeling by alternating back-propagation,

which first infers the latent variable by sampling from its posterior distribution and then updates the model parameters by usual back-propagation. Our SPI shares the same insight with it. To sample efficiently in the continuous latent space, (Tieleman, 2008; Nijkamp et al., 2019) propose the different versions of MCMC to learn the generative model. Specifically, short-run MCMC (Nijkamp et al., 2019) proposes finite-step inference dynamics guided by an energy-based model. We further scale up this idea in the scenarios of PLMs to sample from the continuous latent space.

5. Conclusion

In this work, we propose a probabilistic model with dual latent variables, one discrete latent variable for knowledge selection and one continuous latent variable for response generation. This model is effectively optimized by approximate MLE with the proposed posterior inference method, SPI. Our model has demonstrated its validity and superiority from both theoretical analysis and empirical studies. Further ablation studies show that SPI can search the discrete and continuous space efficiently by our proposed initializer and short-run MCMC in fine-tuning PLMs. Although in this paper, we mainly focus on KGD scenarios, our proposed method, SPI, has the potential to be applied to other knowledge-intensive tasks which require reasoning ability during text generation. We leave further exploration to future work.

References

- Chen, X., Meng, F., Li, P., Chen, F., Xu, S., Xu, B., and Zhou, J. Bridging the gap between prior and posterior knowledge selection for knowledge-grounded dialogue generation. In *Proceedings of the 2020 conference on empirical methods in natural language processing (EMNLP)*, pp. 3426–3437, 2020.
- Cremer, C., Li, X., and Duvenaud, D. Inference suboptimality in variational autoencoders. In *International Conference on Machine Learning*, pp. 1078–1086. PMLR, 2018.
- Dinan, E., Roller, S., Shuster, K., Fan, A., Auli, M., and Weston, J. Wizard of wikipedia: Knowledge-powered conversational agents. In *International Conference on Learning Representations*, 2018.
- Dinan, E., Roller, S., Shuster, K., Fan, A., Auli, M., and Weston, J. Wizard of Wikipedia: Knowledge-powered conversational agents. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2019.
- Durmus, E., He, H., and Diab, M. Feqa: A question answering evaluation framework for faithfulness assessment in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 5055–5070, 2020.
- Feng, S., Patel, S. S., Wan, H., and Joshi, S. Multidoc2dial: Modeling dialogues grounded in multiple documents. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 6162–6176, 2021.
- Fu, T., Zhao, X., Tao, C., Wen, J.-R., and Yan, R. There are a thousand hamlets in a thousand people’s eyes: Enhancing knowledge-grounded dialogue with personal memory. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3901–3913, 2022.
- Ghazvininejad, M., Brockett, C., Chang, M.-W., Dolan, B., Gao, J., Yih, W.-t., and Galley, M. A knowledge-grounded neural conversation model. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32, 2018.
- Han, T., Lu, Y., Zhu, S.-C., and Wu, Y. N. Alternating back-propagation for generator network. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 31, 2017.
- Izacard, G. and Grave, É. Leveraging passage retrieval with generative models for open domain question answering. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pp. 874–880, 2021.
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y., Madotto, A., and Fung, P. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 2022a.
- Ji, Z., Liu, Z., Lee, N., Yu, T., Wilie, B., Zeng, M., and Fung, P. Rho: Reducing hallucination in open-domain dialogues with knowledge grounding. *arXiv preprint arXiv:2212.01588*, 2022b.
- Kim, B., Ahn, J., and Kim, G. Sequential latent knowledge selection for knowledge-grounded dialogue. In *International Conference on Learning Representations*, 2019.
- Kingma, D. P. and Welling, M. Auto-encoding variational bayes. In Bengio, Y. and LeCun, Y. (eds.), *2nd International Conference on Learning Representations, ICLR 2014, Banff, AB, Canada, April 14-16, 2014, Conference Track Proceedings*, 2014. URL <http://arxiv.org/abs/1312.6114>.
- Langevin, P. *On the theory of Brownian motion*. 1908.
- Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., Stoyanov, V., and Zettlemoyer, L. Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. In *ACL 2020*, 2020.
- Li, H. and Han, T. Learning sparse latent representations for generator model. *arXiv preprint arXiv:2209.09949*, 2022.
- Li, J., Galley, M., Brockett, C., Gao, J., and Dolan, W. B. A diversity-promoting objective function for neural conversation models. In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 110–119, 2016.
- Li, L., Xu, C., Wu, W., Zhao, Y., Zhao, X., and Tao, C. Zero-resource knowledge-grounded dialogue generation. *Advances in Neural Information Processing Systems*, 33: 8475–8485, 2020.
- Lian, R., Xie, M., Wang, F., Peng, J., and Wu, H. Learning to select knowledge for response generation in dialog systems.
- Lin, C.-Y. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Liu, S., Zhao, X., Li, B., Ren, F., Zhang, L., and Yin, S. A three-stage learning framework for low-resource knowledge-grounded dialogue generation. *arXiv preprint arXiv:2109.04096*, 2021.

- Meng, C., Ren, P., Chen, Z., Sun, W., Ren, Z., Tu, Z., and Rijke, M. d. Dukenet: A dual knowledge interaction network for knowledge-grounded conversation. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 1151–1160, 2020.
- Moghe, N., Arora, S., Banerjee, S., and Khapra, M. M. Towards exploiting background knowledge for building conversation systems. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 2322–2332, Brussels, Belgium, October–November 2018. Association for Computational Linguistics. doi: 10.18653/v1/D18-1255. URL <https://aclanthology.org/D18-1255>.
- Nijkamp, E., Hill, M., Zhu, S.-C., and Wu, Y. N. Learning non-convergent non-persistent short-run mcmc toward energy-based model. *Advances in Neural Information Processing Systems*, 32, 2019.
- OpenAI. Chatgpt: Optimizing language models for dialogue, Jan 2023. URL <https://openai.com/blog/chatgpt/>.
- Pang, B., Han, T., Nijkamp, E., Zhu, S.-C., and Wu, Y. N. Learning latent space energy-based prior model. In Larochelle, H., Ranzato, M., Hadsell, R., Balcan, M., and Lin, H. (eds.), *Advances in Neural Information Processing Systems*, volume 33, pp. 21994–22008. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/file/fa3060edb66e6ff4507886f9912e1ab9-Paper.pdf>.
- Pang, B., Nijkamp, E., Han, T., and Wu, Y. N. Generative text modeling through short run inference. *arXiv preprint arXiv:2106.02513*, 2021a.
- Pang, B., Zhao, T., Xie, X., and Wu, Y. N. Trajectory prediction with latent belief energy-based model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 11814–11824, June 2021b.
- Papineni, K., Roukos, S., Ward, T., and Zhu, W.-J. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Rashkin, H., Reitter, D., Tomar, G. S., and Das, D. Increasing faithfulness in knowledge-grounded dialogue with controllable features. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 704–718, 2021.
- Robbins, H. and Monro, S. A stochastic approximation method. In *Herbert Robbins Selected Papers*, pp. 102–109. Springer, 1985.
- Roller, S., Dinan, E., Goyal, N., Ju, D., Williamson, M., Liu, Y., Xu, J., Ott, M., Smith, E. M., Boureau, Y.-L., et al. Recipes for building an open-domain chatbot. In *EACL*, 2021.
- Scialom, T., Dray, P.-A., Lamprier, S., Piwowarski, B., Staliano, J., Wang, A., and Gallinari, P. Questeval: Summarization asks for fact-based evaluation. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 6594–6604, 2021.
- Serban, I., Sordoni, A., Bengio, Y., Courville, A., and Pineau, J. Building end-to-end dialogue systems using generative hierarchical neural network models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 30, 2016.
- Shuster, K., Poff, S., Chen, M., Kiela, D., and Weston, J. Retrieval augmentation reduces hallucination in conversation. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pp. 3784–3803, 2021.
- Tieleman, T. Training restricted boltzmann machines using approximations to the likelihood gradient. In *Proceedings of the 25th international conference on Machine learning*, pp. 1064–1071, 2008.
- Xie, J., Zhu, Y., Li, J., and Li, P. A tale of two flows: cooperative learning of langevin flow and normalizing flow toward energy-based model. *arXiv preprint arXiv:2205.06924*, 2022.
- Xu, Y., Ishii, E., Winata, G. I., Lin, Z., Madotto, A., Liu, Z., Xu, P., and Fung, P. Caire in dialdoc21: Data augmentation for information seeking dialogue system. In *Proceedings of the 1st Workshop on Document-grounded Dialogue and Conversational Question Answering (DialDoc 2021)*, pp. 46–51, 2021.
- Xu, Y., Ishii, E., Cahyawijaya, S., Liu, Z., Winata, G. I., Madotto, A., Su, D., and Fung, P. Retrieval-free knowledge-grounded dialogue response generation with adapters. In *Proceedings of the Second DialDoc Workshop on Document-grounded Dialogue and Conversational Question Answering*, pp. 93–107, 2022.
- Yang, C., Lin, Z., Li, J., Meng, F., Wang, W., Wang, L., and Zhou, J. Take: Topic-shift aware knowledge selection for dialogue generation. In *Proceedings of the 29th International Conference on Computational Linguistics*, pp. 253–265, 2022.

- Zhan, H., Shen, L., Chen, H., and Zhang, H. Colv: A collaborative latent variable model for knowledge-grounded dialogue generation. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 2250–2261, 2021.
- Zhang, H., Duckworth, D., Ippolito, D., and Neelakantan, A. Trading off diversity and quality in natural language generation. In *Proceedings of the Workshop on Human Evaluation of NLP Systems (HumEval)*, pp. 25–33, 2021.
- Zhang, Y., Sun, S., Galley, M., Chen, Y.-C., Brockett, C., Gao, X., Gao, J., Liu, J., and Dolan, W. B. Dialogpt: Large-scale generative pre-training for conversational response generation. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pp. 270–278, 2020.
- Zhao, X., Wu, W., Tao, C., Xu, C., Zhao, D., and Yan, R. Low-resource knowledge-grounded dialogue generation. In *International Conference on Learning Representations*, 2019.
- Zhao, X., Wu, W., Xu, C., Tao, C., Zhao, D., and Yan, R. Knowledge-grounded dialogue generation with pre-trained language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 3377–3390, 2020.

A. Theoretical Understanding

In Section 2, we sample from $p_\theta(s, z|C, R)$ approximately. Let $q_\theta(s, z|C, R)$ be the actual distribution of the sampled (s, z) . Given model parameters θ_τ at training iteration τ , the updating rule using the approximate posterior distribution of (s, z) is one-step gradient ascent on the following function,

$$Q(\theta) = \frac{1}{N} \sum_{n=1}^N \mathbb{E}_{q_{\theta_\tau}(s^n, z^n | R^n, C^n)} [\log p_\theta(s^n, z^n, R^n | C^n)]. \quad (23)$$

Comparing to the log-likelihood in Equation (5), we have,

$$\begin{aligned} Q(\theta) &= L(\theta) + \frac{1}{N} \sum_{n=1}^N \mathbb{E}_{q_{\theta_\tau}(s^n, z^n | R^n, C^n)} [\log p_\theta(s^n, z^n | R^n, C^n)] \\ &= L(\theta) - \frac{1}{N} \sum_{n=1}^N \mathbb{D}_{\text{KL}}(q_{\theta_\tau}(s^n, z^n | R^n, C^n) || p_\theta(s^n, z^n | R^n, C^n)) \\ &\quad + \frac{1}{N} \sum_{n=1}^N \mathbb{E}_{q_{\theta_\tau}(s^n, z^n | R^n, C^n)} [\log q_{\theta_\tau}(s^n, z^n | R^n, C^n)]. \end{aligned} \quad (24)$$

With θ_τ fixed, the above equation becomes a function of θ . Then the updating rule follows the stochastic gradient of

$$\tilde{Q}(\theta) = L(\theta) - \frac{1}{N} \sum_{n=1}^N \mathbb{D}_{\text{KL}}(q_{\theta_\tau}(s^n, z^n | R^n, C^n) || p_\theta(s^n, z^n | R^n, C^n)), \quad (25)$$

which can be viewed as a perturbation or variational lower bound of $L(\theta)$.

The fixed point of the learning algorithm that updates θ in Equation (19) solves the following estimating equation:

$$\frac{1}{N} \sum_{n=1}^N \mathbb{E}_{q_{\theta}(s^n, z^n | R^n, C^n)} [\nabla_\theta \log p_\theta(s^n, z^n, R^n | C^n)] = 0. \quad (26)$$

The Monte Carlo approximate of the above expectation becomes the Robbins-Monro algorithm for stochastic approximation (Robbins & Monro, 1985). The convergence to the fixed point follows the regular conditions of the Robbins-Monro algorithm.

B. Loss functions

The gradient update of the model parameter in Equations (19) and (20) is equivalent to minimizing the following total loss function,

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{neg}} + \mathcal{L}_{\text{cls}}, \quad (27)$$

which is a combination of negative log-likelihood loss for Equation (19) and classification loss for Equation (20).

As discussed in Section 2, we treat initializer as a multi-label classifier or multiple binary classifiers,

$$\mathcal{L}_{\text{cls}} = -\frac{1}{N} \sum_{n=1}^N \sum_{i=1}^M y_i^n \log f_\gamma(C_i^n) + (1 - y_i^n) \log(1 - f_\gamma(C_i^n)). \quad (28)$$

where ground-truth $y_i^n = 1$ if it's gold annotation or it's selected by posterior knowledge selection.

For negative log-likelihood loss, with sampled (s^n, z^n) using sequential posterior inference, according to Equation (7),

$$\mathcal{L}_{\text{neg}} = \underbrace{-\frac{1}{N} \sum_{n=1}^N \log p_\theta(z^n | s^n, C^n)}_{\text{prior loss}} - \underbrace{\frac{1}{N} \sum_{n=1}^N \log p_\theta(R^n | s^n, z^n, C^n)}_{\text{reconstruction loss}}. \quad (29)$$

MODEL	WoW SEEN							WoW UNSEEN						
	PPL↓	B3↑	B4↑	R1↑	R2↑	DIST-1↑	DIST-2↑	PPL↓	B3↑	B4↑	R1↑	R2↑	DIST-1↑	DIST-2↑
DRD	23.0	7.5	5.5	18.0	—	—	—	25.6	6.2	4.3	16.5	—	—	—
1/2 DATA	25.3	7.3	5.3	17.5	—	—	—	27.7	6.4	4.5	16.7	—	—	—
1/4 DATA	29.2	6.4	4.4	16.9	—	—	—	32.4	6.0	4.1	16.2	—	—	—
1/8 DATA	33.5	5.9	3.9	16.3	—	—	—	35.8	5.4	3.5	16.0	—	—	—
1/16 DATA	38.6	5.2	3.3	15.7	—	—	—	41.0	5.0	3.2	15.3	—	—	—
KAT-TSLF	14.4	9.1	6.7	21.7	7.6	9.5	38.3	15.8	8.3	6.0	20.7	7.2	6.7	26.0
1/4 DATA	17.6	7.7	5.5	20.3	6.8	9.9	39.1	18.4	7.5	5.2	19.9	6.4	6.6	25.1
1/8 DATA	18.8	7.1	4.9	19.8	6.3	9.9	39.5	20.1	7	4.8	19.0	5.9	6.6	25.3
ZERO DATA	100+	4.0	2.2	14.7	3	7.5	33.9	100+	4.7	2.7	14.9	3	5.7	26.4
SPI (OURS)	17.1	10.2	7.7	22.7	8.8	10.8	40.9	19.1	9.6	7.3	22.0	8.5	6.9	24.3
1/2 DATA	18.2	9.7	7.3	21.8	8.1	10.6	40.6	20.1	9.2	6.9	21.1	7.7	6.5	23.0
1/4 DATA	18.7	9.3	6.9	21.6	7.8	10.1	39.0	20.7	8.9	6.6	20.9	7.3	6.3	23.1
1/8 DATA	20.3	7.9	5.7	20.2	6.7	9.4	35.8	22.0	8.1	6.0	19.6	6.5	5.8	20.7
1/16 DATA	22.0	7.0	4.9	18.7	5.6	8.9	34.0	23.6	7.2	5.2	18.5	5.7	5.7	20.8

Table 7. Automatic evaluation results on WoW test sets under low-resource settings, compared with DRD (Zhao et al., 2019) and KAT-TSLF (Liu et al., 2021). PPL is short for Perplexity; B3 and B4 represent BLEU-3 and BLEU-4; R1 and R2 denote Rouge-1 and Rouge-2; Dist-1 and Dist-2 denote uni-gram and bi-gram distinct metrics.

C. Baseline Models

In this work, we compare the performance of our model with nine other strong baseline models on two KGD benchmarks. We introduce each of them below:

SKT SKT (Kim et al., 2019) is a sequential latent knowledge selection model for multi-turn KGD tasks. Both prior and posterior distribution for knowledge selection are considered sequential processes. The model can keep track of prior and posterior distribution over knowledge, where both distributions are sequentially updated considering the responses in previous turns. We adopt the knowledge candidate selected by SKT to BART for response generation.

FiD Fusion-in-Decoder (FiD) (Izacard & Grave, 2021) is a simple yet effective model for general knowledge-intensive tasks when the context should be augmented with multiple extra documents. The model can be applied to any encoder-decoder-based PLMs. It encodes different context-document pairs in parallel and concatenates all the output hidden states together so that the decoder can attend to all these representations during generation. In this work, we use BART-base as the backbone model for a fair comparison.

DRD DRD (Zhao et al., 2019) propose a disentangled response decoder in order to isolate parameters for generating responses that depend on dialogue context, knowledge inputs, or responses themselves. When generating one token, the decoder needs to do inference with three groups of parameters respectively and decides the final output token with a decoding manager. The proposed model is also pre-trained on the same dialogue corpus as (Liu et al., 2021), thus it is also able to perform KGD generation under a low-resource setting (Table 7).

ZRKGC ZRKGC (Li et al., 2020) focuses on the situation when the real knowledge -grounded dialogue data is not available during the training process. Two latent variables that represent the knowledge for grounding and the rate of grounding are introduced to the model. The generation process is then formalized within a probabilistic framework and optimized via variational inference.

PIPM PIPM is short for “SKT+PIPM+KDBTS” (Chen et al., 2020). The authors propose posterior information prediction and knowledge distillation-based training strategy for knowledge selection. KL divergence is leveraged to bridge the gap between prior and posterior knowledge selection.

DukeNet DukeNet (Meng et al., 2020) explicitly models knowledge tracking and knowledge shifting and formulating their interactions as dual learning without extra external supervision.

	Topic	Chevrolet Corvette
Dialogue History	User	What do you know about the Chevrolet Corvette?
	System	The Chevy Corvette, or "vette" as it is known, is an iconic American sports car that has been produced for half a century.
	User	Do you remember the prince song Little Red Corvette ?
Selected Knowledge	KnowledGPT	Chevrolet Corvette : The first model, a convertible, was introduced at the GM Motorama in 1953 as a concept show car.
	SPI (Ours)	Little Red Corvette: "Little Red Corvette" is a song by American musician Prince.
Response	KnowledGPT	Yes, it was first introduced at the GM Motorama in 1953 as a concept show car.
	SPI (Ours)	I do. It was a song by American musician Prince.

Table 8. One case from test seen set of WoW, comparing the generated response from SPI with that from KnowledGPT.

CoLV CoLV (Zhan et al., 2021) is a Collaborative Latent Variable model, which builds two separate yet collaborative latent spaces to model Similar to our model, it also simultaneously improves the diversity of both knowledge selection and knowledge-aware response generation. However, the model still depends on variational inference for building both latent spaces.

KnowledGPT KnowledGPT (Zhao et al., 2020), as one of the previous SOTA models, equip response generation with a sequential knowledge selector and jointly optimize both the knowledge selector and the response generator with reinforcement learning and curriculum learning. The knowledge selector reranks all the knowledge candidates. The knowledge candidates are concatenated with dialogue history as inputs and truncated to meet the length constraint of the GPT-2 model.

KAT-TSLF KAT-TSLF (Liu et al., 2021) is also one of the previous SOTA models. The authors propose a three-stage learning framework for low-resource knowledge-grounded dialogue tasks. First, the dialogue history encoder and knowledge encoder are pre-trained on the dialogue corpus and knowledge base, respectively. Then, the authors match each dialogue turn in the dialogue corpus with a pseudo gold knowledge from the knowledge base and use the processed new corpus to pre-train the whole model. After two-stage pre-training, the model is adapted to downstream KGD benchmarks and maintains strong performance under low-resource settings. Instead of selecting the knowledge from knowledge candidates, all the provided knowledge sentences are used as inputs, and the decoder is trained to select from all the information.

D. Case Study

Table 8 demonstrate two typical cases from WoW test sets, comparing SPI with KnowledGPT. In the presented case, the dialogue history shows a topic shift from "Chevrolet Corvette" to the "Little Red Corvette" song. However, KnowledGPT fails to capture this shift, whereas our model selects the most relevant knowledge.

E. Human Evaluation

In addition to the automatic evaluation, we conduct the human evaluation to assess the quality of responses generated by our model *SPI* and baseline *KnowledGPT* on WoW. We randomly select 50 samples from each model, and each sample is evaluated by three different annotators. For each comparison, the same context, and two generated responses from each model are shown to the annotators. We require the annotators to be masters with the following qualifications: The numbers of Human Intelligence Tasks (HITs) approved are greater than or equal to 5000, and their HIT approval rates are greater than or equal to 95%. The locations of annotators are restricted to Australia, Canada, the United Kingdom, and the United States. After collecting annotations for AMT, we calculate each score as follows: For A/B testing on Fluency and Relevance, we give one score to the model if it generates an equally good or better response than the other one. We present the ratio of the number of data samples with one score over all the test samples. For 4-point Likert scale on Faithfulness, we assign different levels of faithfulness scores from four to one, then present the average score over all the test samples. Figure 2, Figure 3, and Figure 4 display the annotator instructions of Amazon Mechanical Turk for *Fluency*, *Relevance*, and *Faithfulness*,

Diverse and Faithful Knowledge-Grounded Dialogue Generation via Sequential Posterior Inference

respectively. Please find the instructions and examples for annotators in these figures.

InstructionsShortcuts

Instructions

Please check the reference examples of conversations below.
In this task, you are supposed to determine which response is more fluent.
We also show the topic of the conversation for your reference.
Please see the answers and reasons below.

Definition of Fluency
The response is complete, grammatically correct, and self-consistent without repetition.

REFERENCE EXAMPLE 1:

Topic:
Kale

History:
User: It helps that it's versatile, too, the smoothie bar near me starting serving a kale smoothie, didn't sound great but they add apple and kiwi, pretty tasty.
Assistant: oh really? is your local smoothie bar a planet smoothie or a smoothie king?
User: actually it's tropical cafe & smoothie bar, though i really wish we had a jamba juice.

Model 0 Response:
Assistant: i love jamba juice! i love planet smoothie too! i love the fact that they serve smoothies from the third largest chain of

Model 1 Response:
Assistant: I'm not sure but Jamba Juice was founded in 1990 by Perron.

QUESTION:
Which response is more fluent?

ANSWER:
Model 1 response is more fluent.

Reasons:
Model 0 response is not complete.

[More Instructions](#)

Given the same dialogue context (external knowledge + dialogue history) between the user and assistant, AI assistant generates a dialogue response. We also show the topic of the conversation for your reference.
Please read the context and response and determine which response is more complete, grammatically correct, and self-consistent without repetition. Please also click the Button "Instructions" and read examples, especially the "REFERENCE EXAMPLE" carefully.

Topic:
\${topic}

History:
User: \${human1}
Assistant: \${human2}
User: \${human3}

Model 0 Response:
Assistant: \${model0}

Model 1 Response:
Assistant: \${model1}

Definition of Fluency

The response is complete, grammatically correct, and self-consistent without repetition.

According to the above definition, which response is more fluent?
☐ Both responses are fluent.
☐ Model 0 response is more fluent.
☐ Model 1 response is more fluent.
☐ Neither response is fluent.

Figure 2. The annotator instruction for human evaluation on fluency via A/B testing.

InstructionsShortcuts

Instructions

Please check the reference examples of conversations below.
In this task, you are supposed to determine which response and the corresponding selected knowledge are more relevant to the dialogue history.
We also show the topic of the conversation for your reference.
Please see the answers and reasons below.

Definition of Relevance
The selected knowledge and the corresponding response are relevant to the dialogue history.

REFERENCE EXAMPLE 1:

Topic:
Chevrolet Corvette

History:
User: what do you know about the chevrolet corvette?
Assistant: the chevy corvette, or "vette" as it is known, is an iconic american sports car that has been produced for half a century.
User: do you remember the price song little red corvette?

Knowledge that is selected by model 0:
Chevrolet Corvette: The first model, a convertible, was introduced at the GM Motorama in 1953 as a concept show car.

Model 0 Response:
Assistant: yes, it was first introduced at the gm motorama in 1953 as a concept show car.

Knowledge that is selected by model 1:
Little Red Corvette: Little Red Corvette is a song by the band The Police.

[More Instructions](#)

Given the same dialogue context (external knowledge + dialogue history) between the user and assistant, AI assistant selects the most relevant knowledge sentence and generates a dialogue response. We also show the topic of the conversation for your reference.
Please read the context and response and determine which response and the corresponding selected knowledge are more relevant to the dialogue history. Please also click the Button "Instructions" and read examples, especially the "REFERENCE EXAMPLE" carefully.

Topic:
\${topic}

History:
User: \${human1}
Assistant: \${human2}
User: \${human3}

Knowledge that is selected by model 0:
\${knowledge0}

Model 0 Response:
Assistant: \${model0}

Knowledge that is selected by model 1:
\${knowledge1}

Model 1 Response:
Assistant: \${model1}

Definition of Relevance

The selected knowledge and the corresponding response are relevant to the dialogue history.

According to the above definition, which response and the corresponding selected knowledge are more relevant to the dialogue history?
☐ Both knowledges and responses are relevant.
☐ Model 0 knowledge and response are more relevant.
☐ Model 1 knowledge and response are more relevant.
☐ Neither is relevant.

Figure 3. The annotator instructions for human evaluation on knowledge relevance via A/B testing.

Instructions

Shortcuts

Instructions

Please check the reference examples of conversations below.

In this task, you are supposed to determine whether the response is faithful to the context (external knowledge + history). We also show the topic of the conversation for your reference.

Please see the answers and reasons below.

Definitions and judgement criteria

Fully faithful: The response is fully supported by the dialogue context (external knowledge + history). And the response correctly conveys the information in external knowledge.

Partially faithful: Part of content in the response is supported by the dialogue content. However, some word usage may not be correct.

Not verifiable: The response has no verifiable objective content.

Not faithful: The response is not supported by the dialogue context at all.

External Knowledge consists of two part: [Knowledge topic]:[Knowledge sentence].

REFERENCE EXAMPLE 1:

External Knowledge:

Coors Light: Coors Light has a "mountain icon" to represent the beer in place of the logo.

Topic:

Coors Brewing Company

History:

User: nice, what drinks do they produce?

Assistant: i would say one of there most popular items is "coors light" which was first produced in 1978, so after 105 years of being founded. but they sell all kinds of beer.

User: what share of the market for beer did they capture?

Response:

Assistant: they were the first to make a

[More Instructions](#)

Given the same dialogue context (external knowledge + history) between the user and assistant, AI assistant generates a dialogue response. We also show the topic of the conversation for your reference. Please read the context and response and determine whether each response is faithful to the context. Please also click the Button "Instructions" and read examples, especially the "REFERENCE EXAMPLE" carefully.

External Knowledge:

#{knowledge}

Topic:

#{topic}

History:

User: #{human1}

Assistant: #{human2}

User: #{human3}

Response:

Assistant: #{model}

Definitions and judgement criteria

Fully faithful: The response is fully supported by the dialogue context (external knowledge + history). And the response correctly conveys the information in external knowledge.

Partially faithful: Part of content in the response is supported by the dialogue content. However, some word usage may not be correct.

Not verifiable: The response has no verifiable objective content.

Not faithful: The response is not supported by the dialogue context at all.

According to the above criteria, How faithful is the response?

☐ Faithful
 ☐ Partially faithful
 ☐ Not verifiable
 ☐ Not faithful

Figure 4. The annotator instructions for human evaluation on *Faithfulness* via 4-point Likert scale.

F. Reproducibility

For reproducibility, we provide the code, along with saved checkpoints of our best model with SPI. The evaluation scripts and the instruction for code usage are also included ³.

³https://drive.google.com/drive/folders/1FVRA01uPUVdJ5rzN_mFcwpZAznr8XDIX?usp=share_link